

# CMPSC 240A HW-3 Report

Tanmoy Sanyal

October 23, 2016

## Brief code summary

The code follows the logic outlined in the description of the assignment. `rec_cilkified` works by divide and conquer. The arrays are divided into two (not always equal) parts recursively and their individual dot products are summed up, until a certain threshold (COARSENESS) is reached. After that point, the dot products are done in a simple serial for loop. `cilk_spawn` is used in the splitting so that the recursive function calls are spawned off in different threads. `loop_cilkified` works in a similar way, but instead of recursion, it splits the arrays into chunks of length = COARSENESS and then these chunks are "dot"-ed individually and added up. `hyperobject_cilkified` simply uses the hyperobject construct of cilk++ and it takes care of everything.

For all the scaling experiments, I have assumed that the number of threads is the same as the number of cores (and set using the CILK\_NWORKERS environment variable) i.e. hyperthreading has been assumed to be absent.

## Scaling with size of array

For this experiment, the COARSENESS was chosen to be 50, and was run on all 24 cores of a single Comet compute node. The array sizes were varied from  $10^4$  to  $10^{10}$ . However, the sequential algorithm run-time for sizes below  $10^6$  is almost negligible, and reported as 0.

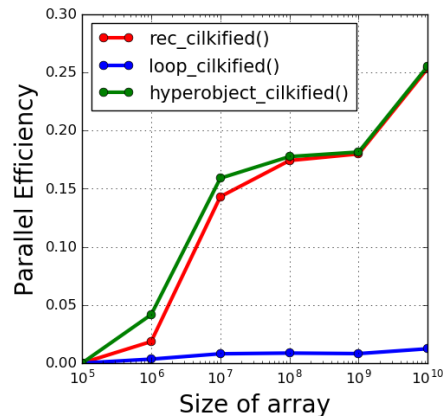


Figure 1: Parallel efficiency vs array size, NCores = 24, COARSENESS = 50

## Scaling with number of cores

For this experiment, the array size was fixed at  $10^6$  and COARSENESS was maintained at 50. The number of cores was varied from 1 to 64 in powers of 2. Note that the work ( $t_1$ ) for this experiment is not the runtime of the sequential function `std::inner_product`, but the time obtained with 1 core and `CILK_NWORKERS = 1`. (For 32 and 64 cores, I use 16 cores per compute node and multiple nodes).

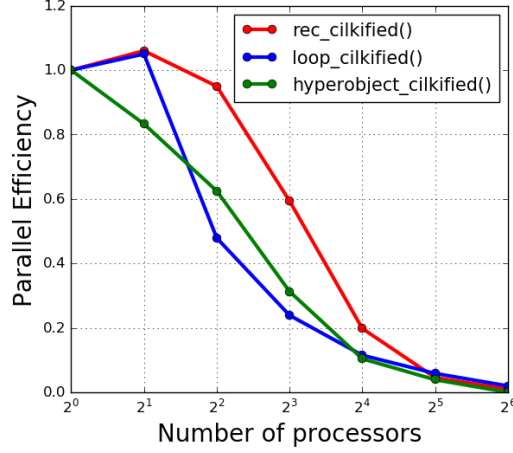


Figure 2: Parallel efficiency vs number of processors, Array size =  $10^6$ , COARSENESS = 50

## Sensitivity to COARSENESS

For this experiment, the array size was fixed at  $10^8$  and was run on all 24 cores of a single Comet compute node. COARSENESS was varied from 1 to  $10^8$  in powers of 10.

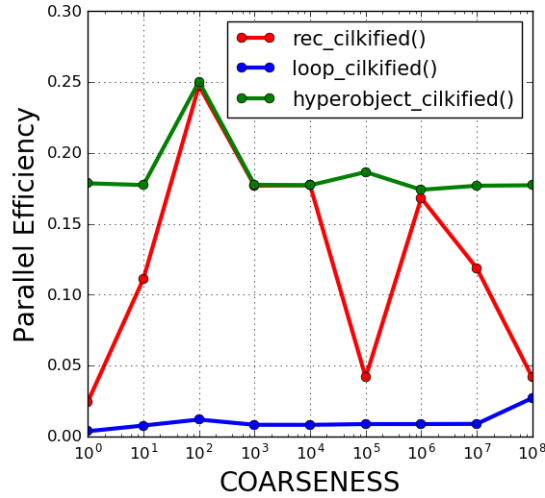


Figure 3: Parallel efficiency vs COARSENESS, Array size =  $10^8$ , NCores = 24

## Best performing dot-product code

From the different scaling studies presented above, Fig. 1 shows that with increasing number of cores (and therefore spawned threads), the performance goes down. Thus, the performance is never truly a function of how much computational resources I have at my disposal. On the other hand, increasing array sizes is a very real world phenomenon, and it forms the primary metric on which to judge performance. For large array sizes, `rec_cilkified` and `hyperobject_cilkified` have nearly similar parallel efficiency in Fig. 2. So any of these two might be a good candidate. Now, coming to COARSENESS, we see from Fig. 3, that `rec_cilkified` is much more sensitive than `hyperobject_cilkified`. That means that the former is amenable to easy tuning.

On the basis of the above insights, I would use the following parameters for the best-performing dot product code:

- Function: `rec_cilkified()`
- COARSENESS: 100

The speedup obtained using these settings is: Clearly, this does not enjoy linear speedup. Beyond

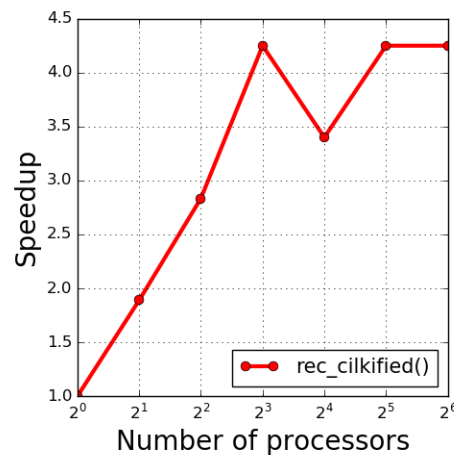


Figure 4: Speedup vs number of cores with `rec_cilkified`, Array size =  $10^6$ , COARSENESS = 100

16 processors, the speedup saturates, because the threads are no longer in truly shared memory and span different Comet nodes. That is a trivial bottleneck. But the speedup also falls (and drastically so) from 8 to 16 processors, which are all within the same node. TODO: Reason

## Tracing parallelism in the code